
Forecasting Energy Consumption - Kansas City, Kansas

Project Summary:

- This project focuses on forecasting energy consumption for Kansas City, Kansas. Data from the Energy Information Administration (EIA) and weather data from Visual Crossing are used. The goal is to build a machine learning model capable of predicting energy use specifically for this region, using historical data from 2020 to 2024.
- The dataset combines hourly energy consumption records and weather data. Key features such as hour, day of the week, month, and specific holidays, are created to capture time-based patterns. Temperature is included to account for weather-related variability.
- An XGBoost regression model is trained and evaluated using time-series cross-validation, achieving an average RMSE of approximately 5.84% of the energy consumption range. The model is then used to predict future energy consumption for Kansas City, Kansas. This tool aims to support energy management and utility planning for the region.

```
In [1]: import holidays

import requests
import json

import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('fivethirtyeight')
color_pal = sns.color_palette()

import xgboost as xgb
```

```
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import TimeSeriesSplit, GridSearchCV
from sklearn.utils import resample
```

API Call for Energy Data (eia.gov)

```
In [2]: with open('eia_api_key') as f:
        api_key = f.read()
```

```
In [3]: base_url = ('https://api.eia.gov/v2/electricity/rto/region-sub-ba-data/data/'
                    '?frequency=hourly&data[0]=value&facets[subba][]=KACY&start=2020-01-01T00'
                    '&end=2024-12-28T22&sort[0][column]=period&sort[0][direction]=asc'
                    f'&length=5000&api_key={api_key}')
```

```
In [4]: offset = 0
        all_data = []

        while True:
            url = f"{base_url}&offset={offset}"
            response = requests.get(url)
            if response.status_code == 200:
                data = response.json()['response']['data']
                if not data:
                    break
                all_data.extend(data)
                offset += 5000
            else:
                print(f"Request failed with status code: {response.status_code}")
                print(response.text)
                break
```

make dataframe

```
In [5]: df = pd.DataFrame(all_data)
```

```
In [6]: df.head(3)
```

```
Out[6]:
```

	period	subba	subba-name	parent	parent-name	value	value-units
0	2020-01-01T00	KACY	Kansas City Board of Public Utilities - SWPP	SWPP	Southwest Power Pool	259	megawatthours
1	2020-01-01T01	KACY	Kansas City Board of Public Utilities - SWPP	SWPP	Southwest Power Pool	263	megawatthours
2	2020-01-01T02	KACY	Kansas City Board of Public Utilities - SWPP	SWPP	Southwest Power Pool	258	megawatthours

```
In [7]: df['value'] = df.value.astype(int)
df['period'] = pd.to_datetime(df['period'])
```

```
In [8]: df = df[['period', 'value']]
df.rename(columns={'period': 'datetime'}, inplace=True)
df.head()
```

```
Out[8]:
```

	datetime	value
0	2020-01-01 00:00:00	259
1	2020-01-01 01:00:00	263
2	2020-01-01 02:00:00	258
3	2020-01-01 03:00:00	253
4	2020-01-01 04:00:00	247

```
In [9]: print(f'Missing: {df.isna().sum().sum()}')
print(f'Duplicates: {df.index.duplicated().sum()}')
```

Missing: 0
Duplicates: 0

Hourly Weather Data (visualcrossing.com)

```
In [10]: df2 = pd.read_csv(r'C:\Users\baile\Downloads\kansas city, kansas 2020-01-01 to 2024-12-28.csv')
```

```
In [11]: df2.head(3)
```

```
Out[11]:
```

	name	datetime	temp	feelslike	dew	humidity	precip	precipprob	preciptype	snow	...	sealevelpressure	cloudcover
0	kansas city, kansas	2020-01-01T00:00:00	-2.1	-6.2	-6.1	74.00	0.0	0	NaN	0.0	...	1014.5	0.0
1	kansas city, kansas	2020-01-01T01:00:00	-0.7	-3.4	-5.7	68.95	0.0	0	NaN	0.0	...	1013.6	0.0
2	kansas city, kansas	2020-01-01T02:00:00	-1.2	-1.2	-6.2	68.87	0.0	0	NaN	0.0	...	1012.5	0.0

3 rows × 24 columns



```
In [12]: df2['datetime'] = pd.to_datetime(df2['datetime'])
df2 = df2[['datetime', 'temp']]
```

```
In [13]: df2.head()
```

```
Out[13]:
```

	datetime	temp
0	2020-01-01 00:00:00	-2.1
1	2020-01-01 01:00:00	-0.7
2	2020-01-01 02:00:00	-1.2
3	2020-01-01 03:00:00	-0.7
4	2020-01-01 04:00:00	-0.1

```
In [14]: print(f'Missing: {df2.isna().sum().sum()}')
print(f'Duplicates: {df2.index.duplicated().sum()}')
```

Missing: 0

Duplicates: 0

Merge Energy and Weather Data on Datetime

```
In [15]: df = pd.merge(df, df2, on='datetime', how='inner')
```

```
In [16]: df = df.set_index('datetime')
```

```
In [17]: df.head()
```

```
Out[17]:
```

	value	temp
datetime		
2020-01-01 00:00:00	259	-2.1
2020-01-01 01:00:00	263	-0.7
2020-01-01 02:00:00	258	-1.2
2020-01-01 03:00:00	253	-0.7
2020-01-01 04:00:00	247	-0.1


```
In [18]: print(f'Missing: {df.isna().sum().sum()}')  
print(f'Duplicates: {df.index.duplicated().sum()}')
```

```
Missing: 0  
Duplicates: 5
```

```
In [19]: df = df[~df.index.duplicated(keep='first')]
```

Visualizing MWh Usage


```
In [20]: sns.boxplot(df.value)  
plt.title('Boxplot of MWh');
```




- This value being so extreme makes it seem like an error, so it will be dropped.

```
In [21]: df = df[df.value<30000]
```

```
In [22]: sns.boxplot(df.value)  
plt.title('Boxplot of MWh');
```




```
In [23]: df.value.plot(style='.', figsize=(10,5), color=color_pal[0], title='Kansas City Energy Use in MWh')
plt.xticks(rotation=45);
```



Train/Test Split

```
In [24]: train = df.loc[df.index < '01-01-2024']
test = df.loc[df.index >= '01-01-2024']
```

```
In [25]: fig, ax = plt.subplots(figsize=(10,5))
train.value.plot(ax=ax, label='Training Set', title='Data Train/Test Split')
test.value.plot(ax=ax, label='Test Set')
ax.axvline('01-01-2024', color='black', ls='--')
ax.legend(['Training Set', 'Test Set'], loc='upper right', bbox_to_anchor=(1, 1.1));
```



- The test set is the last 362 days of the data.

Feature Creation

```
In [26]: def create_features(df):  
    '''create times series features'''  
    df['hour'] = df.index.hour  
    df['dayofweek'] = df.index.day_of_week  
    df['quarter'] = df.index.quarter  
    df['month'] = df.index.month  
    df['year'] = df.index.year  
    df['dayofyear'] = df.index.dayofyear  
  
    us_holidays = holidays.US()  
    specific_holidays = {"New Year's Day", "Labor Day", "Thanksgiving", "Christmas Day"}  
    df['is_specific_holiday'] = df.index.map(lambda x: us_holidays.get(x.date()) in specific_holidays)
```



```
    return df  
  
create_features(df);
```

- Creating features such as hour and is_specific_holiday should make for a better model.

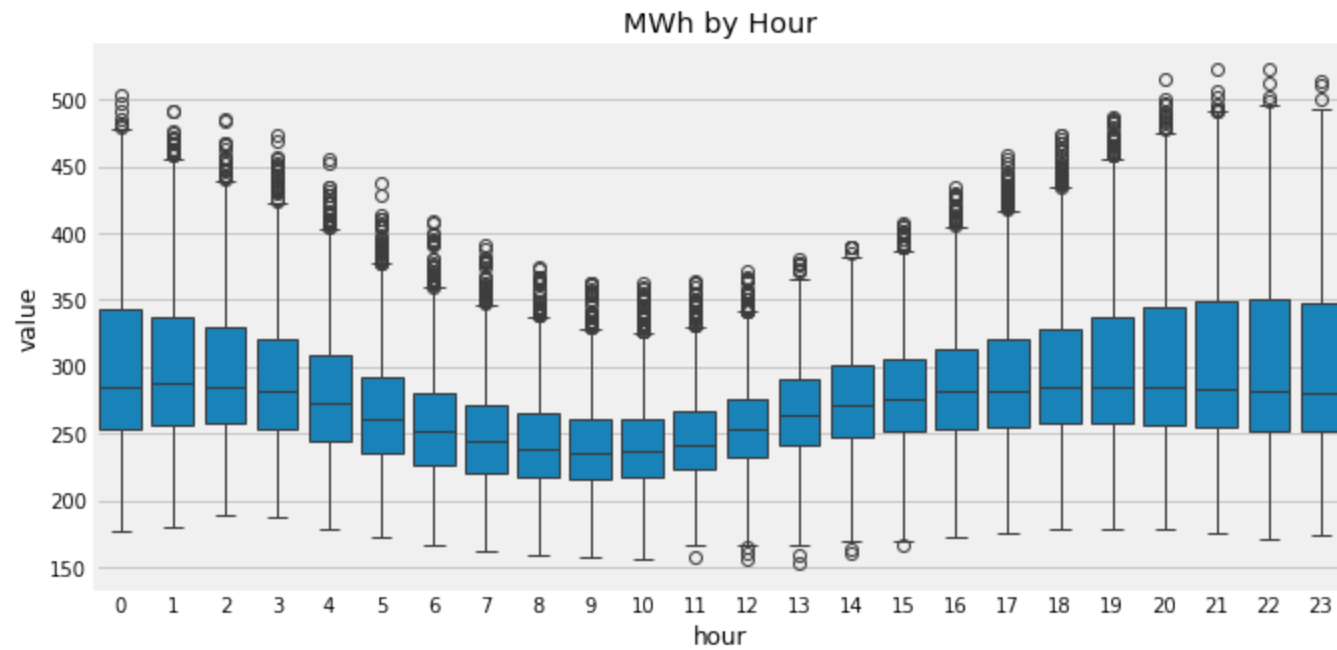
```
In [27]: df.head(3)
```

```
Out[27]:
```

	value	temp	hour	dayofweek	quarter	month	year	dayofyear	is_specific_holiday
datetime									
2020-01-01 00:00:00	259	-2.1	0	2	1	1	2020	1	True
2020-01-01 01:00:00	263	-0.7	1	2	1	1	2020	1	True
2020-01-01 02:00:00	258	-1.2	2	2	1	1	2020	1	True

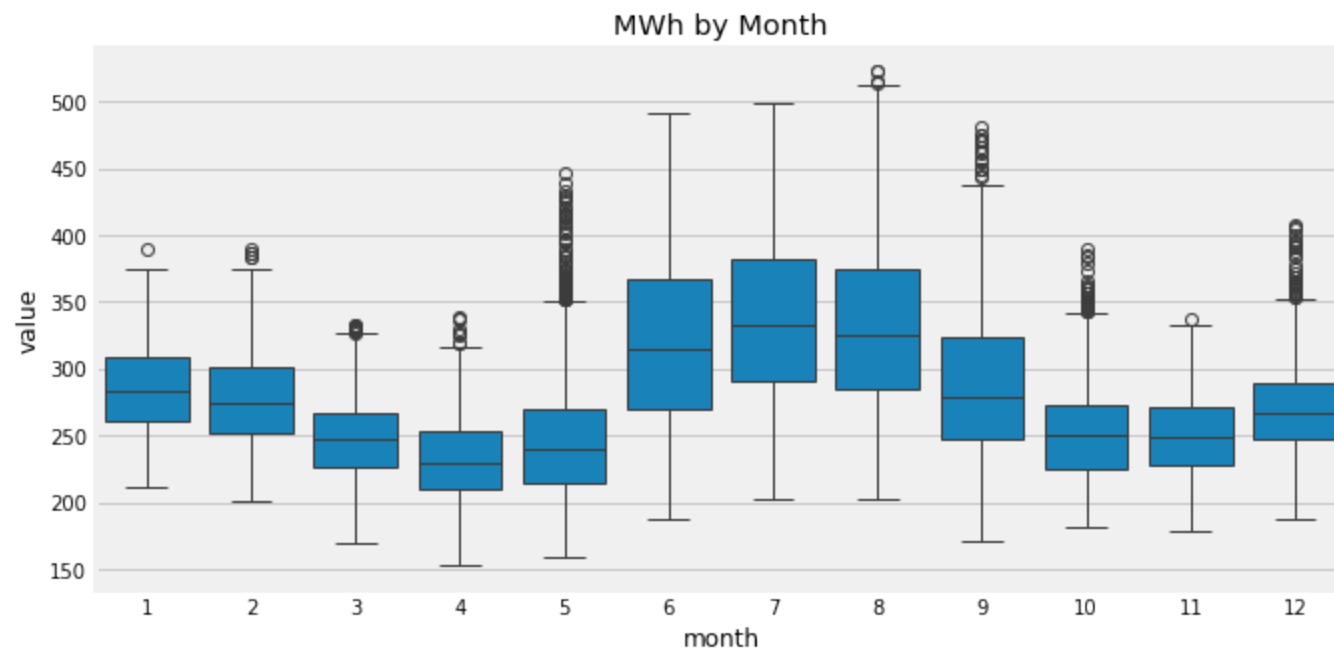
Visualizing Feature/Target Relationships

```
In [28]: plt.figure(figsize=(10,5))  
sns.boxplot(data=df, x='hour', y='value')  
plt.title('MWh by Hour');
```



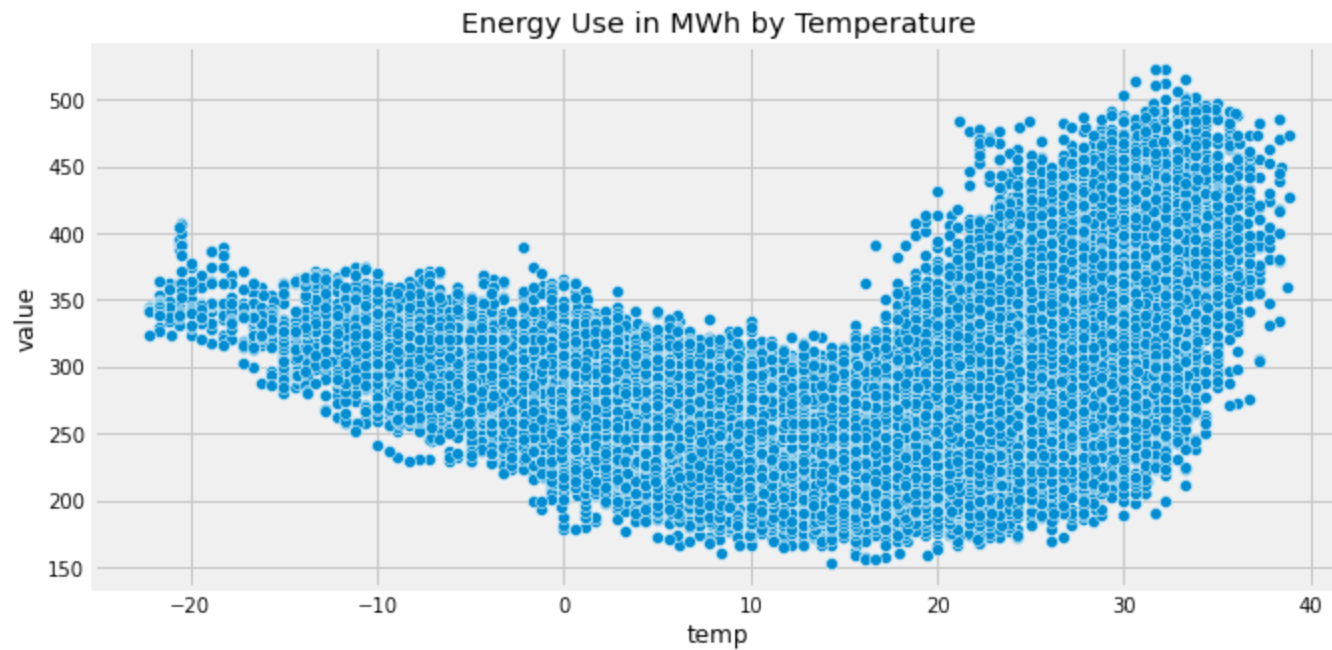
-
- Energy consumption appears to be lowest from about 8 to 10 in the morning, and highest from about 8 to 11 at night.
-

```
In [29]: plt.figure(figsize=(10,5))
sns.boxplot(data=df, x='month',y='value')
plt.title('MWh by Month');
```




- The peaks in energy consumption during summer and the dips in fall align with changes in temperature.

```
In [30]: plt.figure(figsize=(10,5))
sns.scatterplot(x=df.temp, y=df.value)
plt.title('Energy Use in MWh by Temperature');
```



-
- Energy usage appears to be lowest in mild temperatures and highest in hot temperatures. The reason for this graph not being u-shaped could be explained by natural gas being widely used in Kansas City during the winter.
-

```
In [31]: plt.figure(figsize=(10,5))
sns.barplot(data=df, x='is_specific_holiday', y='value', errorbar='sd')
plt.xlabel('Category')
plt.ylabel('Mean Value')
plt.title('Comparison of Means (Holiday vs. Non-Holiday)');
```



- Mean energy usage does not appear to be higher on holidays. That's probably because the specified holidays: New Year's Day, Labor Day, Thanksgiving, Christmas Day, are around winter.

XGBoost Model

```
In [32]: train = create_features(train.copy())
test = create_features(test.copy())

features = ['hour', 'dayofweek', 'quarter', 'month', 'year', 'dayofyear', 'temp', 'is_specific_holiday']
target = 'value'

X_train = train[features]
y_train = train[target]
```

```
X_test = test[features]
y_test = test[target]
```

```
In [33]: reg = xgb.XGBRegressor(n_estimators=1000, early_stopping_rounds=50, random_state=1)
reg.fit(X_train,y_train,
        eval_set=[(X_train,y_train), (X_test,y_test)],
        verbose=100)
```

```
[0]    validation_0-rmse:198.72407    validation_1-rmse:202.91748
[67]    validation_0-rmse:12.34317    validation_1-rmse:22.76310
```


```
Out[33]: XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, early_stopping_rounds=50,
             enable_categorical=False, eval_metric=None, feature_types=None,
             gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
             interaction_constraints=None, learning_rate=None, max_bin=None,
             max_cat_threshold=None, max_cat_to_onehot=None,
             max_delta_step=None, max_depth=None, max_leaves=None,
             min_child_weight=None, missing=nan, monotone_constraints=None,
```

-
- Early stopping terminated the training process after 68 iterations. This indicates that the validation set error stopped improving further. Additional iterations would lead to overfitting on the training data.
-

Feature Importance

```
In [34]: importance_df = pd.DataFrame(data=reg.feature_importances_,
                                     index=reg.feature_names_in_,
                                     columns=['importance'] )
```

```
In [35]: importance_df.sort_values('importance').plot(kind='barh',title='Feature Importance');
```




- Unsurprisingly, the temperature data turned out to be the most important feature.


Forecast Made on the Test Set

```
In [36]: test['prediction'] = reg.predict(X_test)
df = df.merge(test[['prediction']], how='left', left_index=True, right_index=True)
```

```
In [37]: ax = df[['value']].plot(figsize=(10,5))
df['prediction'].plot(ax=ax, style='.')
plt.legend(['True Data', 'Predictions'])
ax.set_title('Raw Data and Predictions');
```



```
In [38]: ax = df.loc[(df.index>'01-01-2024') & (df.index<'12-28-2024')]['value']\  
        .plot(figsize=(10,5), title='Predictions over Test Set')  
  
        df.loc[(df.index>'01-01-2024') & (df.index<'12-28-2024')]['prediction'].plot(style='.')  
  
        plt.legend(['True Data', 'Prediction']);
```

- Visually, the fit looks pretty good. A quantitative measure like rmse would be more informative though.

Time Series Cross Val & Test Score

```
In [39]: X = df.drop(columns=['value', 'prediction'])  
         y = df.value
```

```
In [40]: tscv = TimeSeriesSplit(n_splits=5)  
  
         rmse_scores = []  
  
         for train_index, test_index in tscv.split(X):  
             X_train, X_test = X.iloc[train_index], X.iloc[test_index]  
             y_train, y_test = y.iloc[train_index], y.iloc[test_index]
```

```

reg.fit(
    X_train, y_train,
    eval_set=[(X_train, y_train), (X_test, y_test)], verbose=0)

y_pred = reg.predict(X_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
rmse_scores.append(rmse)

rmse_scores_rounded = [round(score, 2) for score in rmse_scores]
print("RMSE scores for each fold:", rmse_scores_rounded)
print("Average RMSE:", np.mean(rmse_scores))

```

RMSE scores for each fold: [24.93, 26.8, 22.73, 24.25, 21.6]
Average RMSE: 24.061008290912728

-
- The cross val scores are close enough together to indicate that the model's predictions are stable across different subsets of the data.
-

```

In [56]: final_model = xgb.XGBRegressor(n_estimators=1000, early_stopping_rounds=50, random_state=1)
final_model.fit(X_train, y_train, eval_set=[(X_train, y_train), (X_test, y_test)], verbose=0)

y_pred_test = final_model.predict(X_test)
test_rmse = np.sqrt(mean_squared_error(y_test, y_pred_test))
print(f'The test RMSE is {test_rmse.round(2)}.')

```

The test RMSE is 21.6.

```

In [42]: my_range = df.value.max()-df.value.min()
avg_error_perc = round((test_rmse/my_range)*100,2)

print(f"The model's performance on the test set indicates an average error of {avg_error_perc}%.")

```

The model's performance on the test set indicates an average error of 5.84%.

Future Dataframe

-
- The future dataframe will be temperature predictions in Kansas City for the next 16 days along with the created features used earlier.
-

```
In [43]: future_df = pd.read_csv(r'C:\Users\baile\Downloads\kansas city, kansas 2024-12-28 to 2025-12-28 (3).csv')
```

```
In [44]: future_df = future_df[['datetime', 'temp']]
future_df['datetime'] = pd.to_datetime(future_df['datetime'])
future_df = future_df.set_index('datetime')
future_df.head()
```

Out[44]:

	temp
datetime	
2024-12-28 00:00:00	5.0
2024-12-28 01:00:00	4.4
2024-12-28 02:00:00	6.1
2024-12-28 03:00:00	7.2
2024-12-28 04:00:00	7.2

datetime	
2024-12-28 00:00:00	5.0
2024-12-28 01:00:00	4.4
2024-12-28 02:00:00	6.1
2024-12-28 03:00:00	7.2
2024-12-28 04:00:00	7.2

```
In [45]: future_df = create_features(future_df)
```

```
In [46]: future_df.head()
```

Out[46]:

	temp	hour	dayofweek	quarter	month	year	dayofyear	is_specific_holiday
datetime								
2024-12-28 00:00:00	5.0	0	5	4	12	2024	363	False
2024-12-28 01:00:00	4.4	1	5	4	12	2024	363	False
2024-12-28 02:00:00	6.1	2	5	4	12	2024	363	False
2024-12-28 03:00:00	7.2	3	5	4	12	2024	363	False
2024-12-28 04:00:00	7.2	4	5	4	12	2024	363	False

Predicting Future Energy Consumption(2024-01-02 to 2024-02-02)

bootstrap resampling for confidence intervals

```
In [47]: future_predictions = []
for _ in range(100):
    X_boot, y_boot = resample(X_train, y_train)

    reg = xgb.XGBRegressor(n_estimators=68)
    reg.fit(X_boot, y_boot, verbose=0)

    pred = reg.predict(future_df)
    future_predictions.append(pred)
```

```
In [48]: future_predictions = np.array(future_predictions)
lower_bound = np.percentile(future_predictions, 2.5, axis=0)
upper_bound = np.percentile(future_predictions, 97.5, axis=0)
```


plot

```
In [49]: future_df['pred'] = reg.predict(future_df)

plt.figure(figsize=(10, 5))

plt.plot(future_df.index, future_df['pred'], color=color_pal[4], ms=1, lw=1, label='Predicted')
plt.fill_between(future_df.index, lower_bound, upper_bound, color='gray', alpha=0.2, label='95% Confidence Interval')

plt.title('Future Predictions with Confidence Intervals')
plt.xlabel('Datetime')
plt.ylabel('Energy Consumption (MWh)')
plt.xticks(rotation=45)
plt.legend();
```



-
- The plot above shows the predicted energy use in MWh in Kansas City from 2024-12-28 to 2025-01-12.
-

predictions dataframe

```
In [50]: mean_predictions = future_predictions.mean(axis=0)
```

```
predictions_df = pd.DataFrame({
    'datetime': future_df.index,
    'pred': mean_predictions,
    'lower_bound': lower_bound,
    'upper_bound': upper_bound })

predictions_df.set_index('datetime', inplace=True)
```

```
In [51]: predictions_df.head()
```

```
Out[51]:
```

	pred	lower_bound	upper_bound
--	------	-------------	-------------

datetime			
2024-12-28 00:00:00	261.541901	244.706550	274.953561
2024-12-28 01:00:00	262.395569	245.559415	276.466607
2024-12-28 02:00:00	259.488403	248.880390	269.882551
2024-12-28 03:00:00	255.511734	245.261096	265.157507
2024-12-28 04:00:00	244.360703	234.923753	253.269280

```
predictions_df.to_csv('kc_energy_preds_2024_12_28_to_2025_01_12.csv')
```